

Version Control with Git

CSE 3311

Christoph Csallner

University of Texas at Arlington (UTA)

These slides contain material from the sources listed on the following slides titled “Sources”

Sources

- [CD] “What’s Wrong with Git? A Conceptual Design Analysis”
 - By Santiago Perez De Rosso and Daniel Jackson
 - <http://dl.acm.org/citation.cfm?id=2509584>
- [Mike] “Version control concepts and best practices”
 - By Mike Ernst
 - <http://homes.cs.washington.edu/~mernst/advice/version-control.html#dvc-advice>

Why Version Control?

- Manage multiple versions of a piece of information
- Many non-developers do this manually
 - Every time you modify a file: Save file under a new name
 - For example: Foo-1.txt → Foo-2.txt → Foo-3.txt
 - [+] Easy initially
 - [+] Do not need (to learn) any tools besides OS
 - [-] Waste of disk space (redundant copies of unchanged parts)
 - [-] Cumbersome if file name expected to remain constant (e.g., Java source files)
 - [-] Sharing files across machines cumbersome (email?)
 - [-] No support for team work (parallel edits & merge?)

Version Control via Cloud Storage

- Cloud storage mostly good for non-developers
- Every time you modify a file: Save (same name)
 - For example: Foo.txt → Foo.txt → Foo.txt
- Storage client stores new version in cloud
 - Some providers store last N versions of each file
 - [+] Some clients allow reverting to older versions
 - [+] Relatively simple clients + web-based administration, e.g., Microsoft OneDrive, Google Drive, DropBox, Box
 - [+] Little disk space wasted locally (but maybe in cloud)
 - [+] Sharing files across machines easy (auto-sync)
 - [+] Some support for teams (parallel changes to same doc)

Is Cloud Storage Sufficient?

- Cloud Storage seems to do a lot of good things
- But limited support for things developers also do
 - Name and annotate each change
 - Refer to change from other tools (e.g., issue tracking)
 - Many parallel branches of same files
 - Fork and merge branches
 - Merge conflicting changes in same branch

Version Control

- Many alternative names
 - Revision control
 - Configuration management
 - Software configuration management
 - Source code management
 - Source code control
 - Source control
- Essentially all refer to very similar concepts
 - Other sources may provide contradicting definitions

Concrete Example Tools

- Following are a few representative example tools
 - There are many others
http://en.wikipedia.org/wiki/Revision_control
- Proprietary
 - 1970s: client-server
 - 1990s: Bitkeeper (distributed)
- Open-source
 - 1990: CVS (client-server)
 - 2000: Subversion (client-server)
 - 2005: Git, Mercurial=Hg (distributed)

Version Control Benefits

- New version does not require new file name
- Keeps track of project history: For every change
 - What changed, who changed it when
 - Why it was made (if change is annotated)
 - Change can be referenced from external tools
- Helps recover from mistakes: Revert to old version
- Supports collaboration
 - Easy sharing files across machines
 - Create and merge many branches of same files
 - Identify and merge conflicting changes in same branch

Benefits Grow with Team Size

- Assume 500 software developers working on the same project ..
- .. without any kind of version control tools
- The project is doomed
- Given the growing benefits, larger teams are more willing to use and invest in version control tools
- Even if these tools
 - Are hard to learn
 - Have known problems (scalability, etc.)

But Also Scales Down to Single User

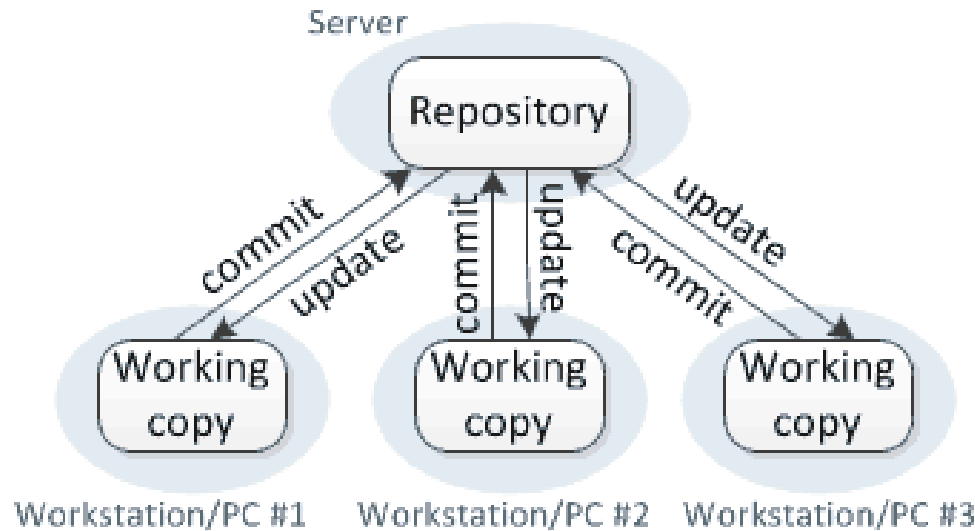
- Very useful even for lone user
- Cloud Storage
 - Syncs files across different machines of single person
 - Creates backups of files (one per machine + cloud)
 - Allows reverting to old versions
- Version Control
 - Similar benefits
 - Can work on multiple branches of same file
 - Easy to compare and merge different versions

Three Tool Generations


1. Single computer
2. Client-server (e.g.: Subversion)
 - Typically server is only node storing full version history
 - [+] Several people can work on same project
 - [-] Single-server bottleneck
3. Peer-to-peer (“distributed”) (e.g.: Git, Mercurial)
 - Each peer has its own client and its own server with full version history → Enables many interaction styles
 - [+] No single bottleneck if used peer-to-peer style
 - [+] Can simulate client-server style (GitHub, BitBucket)
 - [-] May use a lot of space on each peer (i.e., for binaries)

Comparison in Centralized Scenario

Centralized version control



Distributed version control



Following slides are based on [CD]
(except for domain model refresher slide)

GIT

In Git: “File” means “Path of the file”

- “File”
 - Git’s heuristic of the identity of an actual file in your OS’s file system (\neq file contents)
 - In Git: “File” means **path of the file**
 - “File path” would have been a less confusing name
 - Git does not keep track of the identity of an actual file independently of the file path
- Having “file” a/b/c in Git does not mean you have a/b/c in your machine’s file system
 - For example, in Git: “File” a/b/c may be marked as “staged for removal”

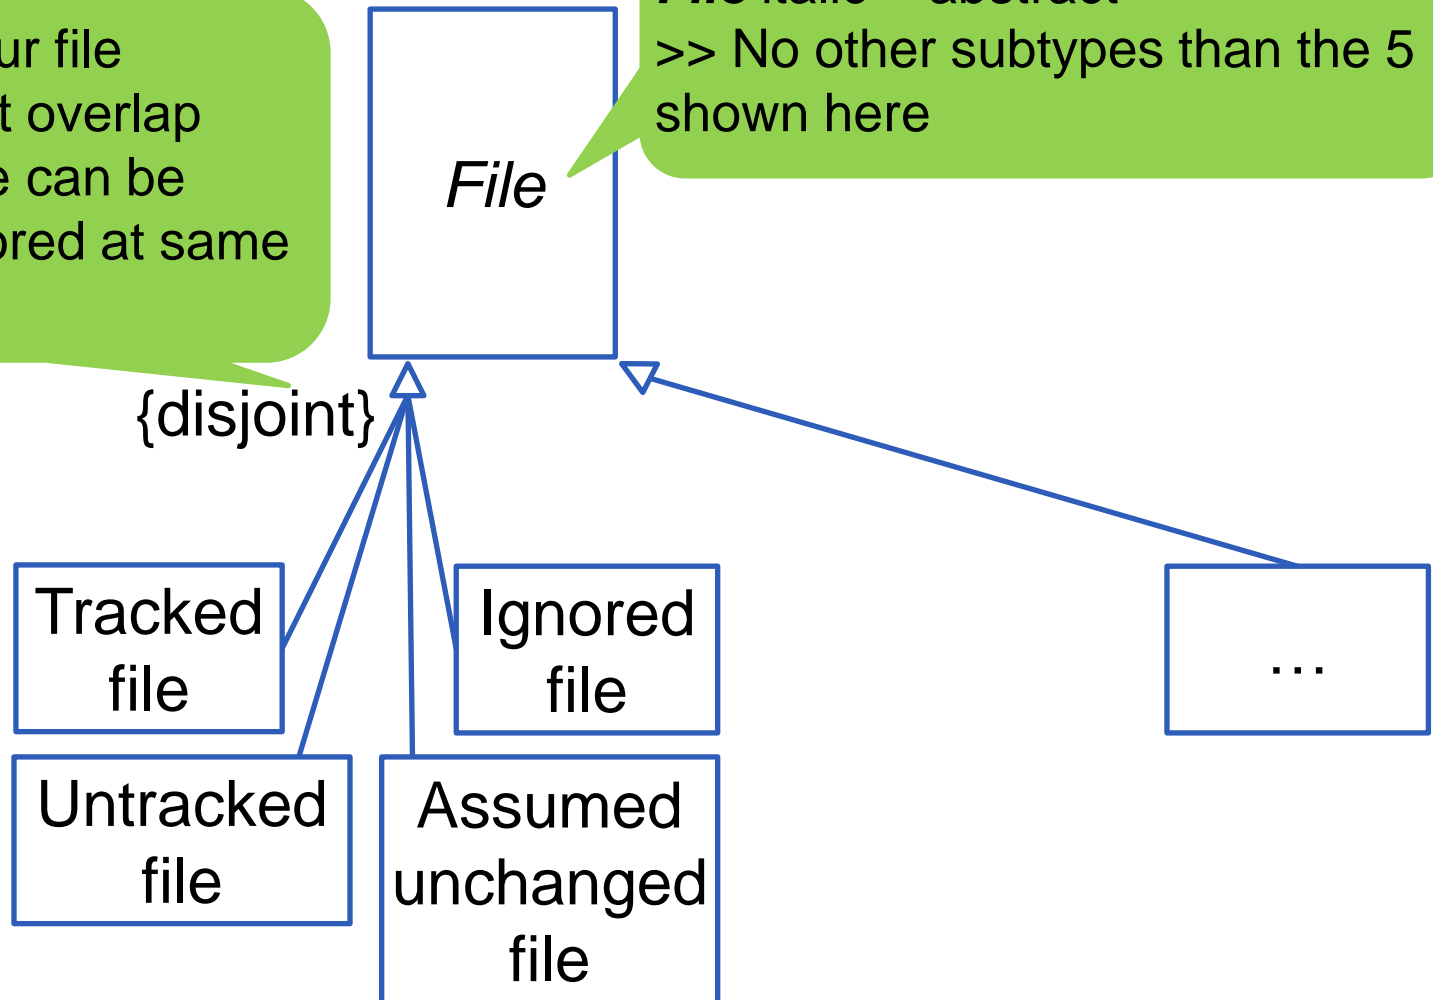
Rename / move an actual file → Git ?

- Aside: Unlike Git, OS keeps track of actual file's identity separately from actual file's name and path
 - Rename actual file: OS only updates name
 - Move actual file: OS only updates path
- Breaks Git's heuristic of the actual file's identity
 - In Git: Delete a/b/c then add a/b/d (or add a/e/c)
 - “Git has a rename command `git mv`, but that is just for convenience. **The effect is indistinguishable from removing the file and adding another with different name and the same content.**”
 - [https://git.wiki.kernel.org/index.php/GitFaq#Why_does_Git_not_.22track.22_renames.3F]
 - Breaks commit history of actual file

Conceptual model: Core file subtypes

{disjoint} = Four file subtypes don't overlap
>> E.g., no file can be tracked & ignored at same time

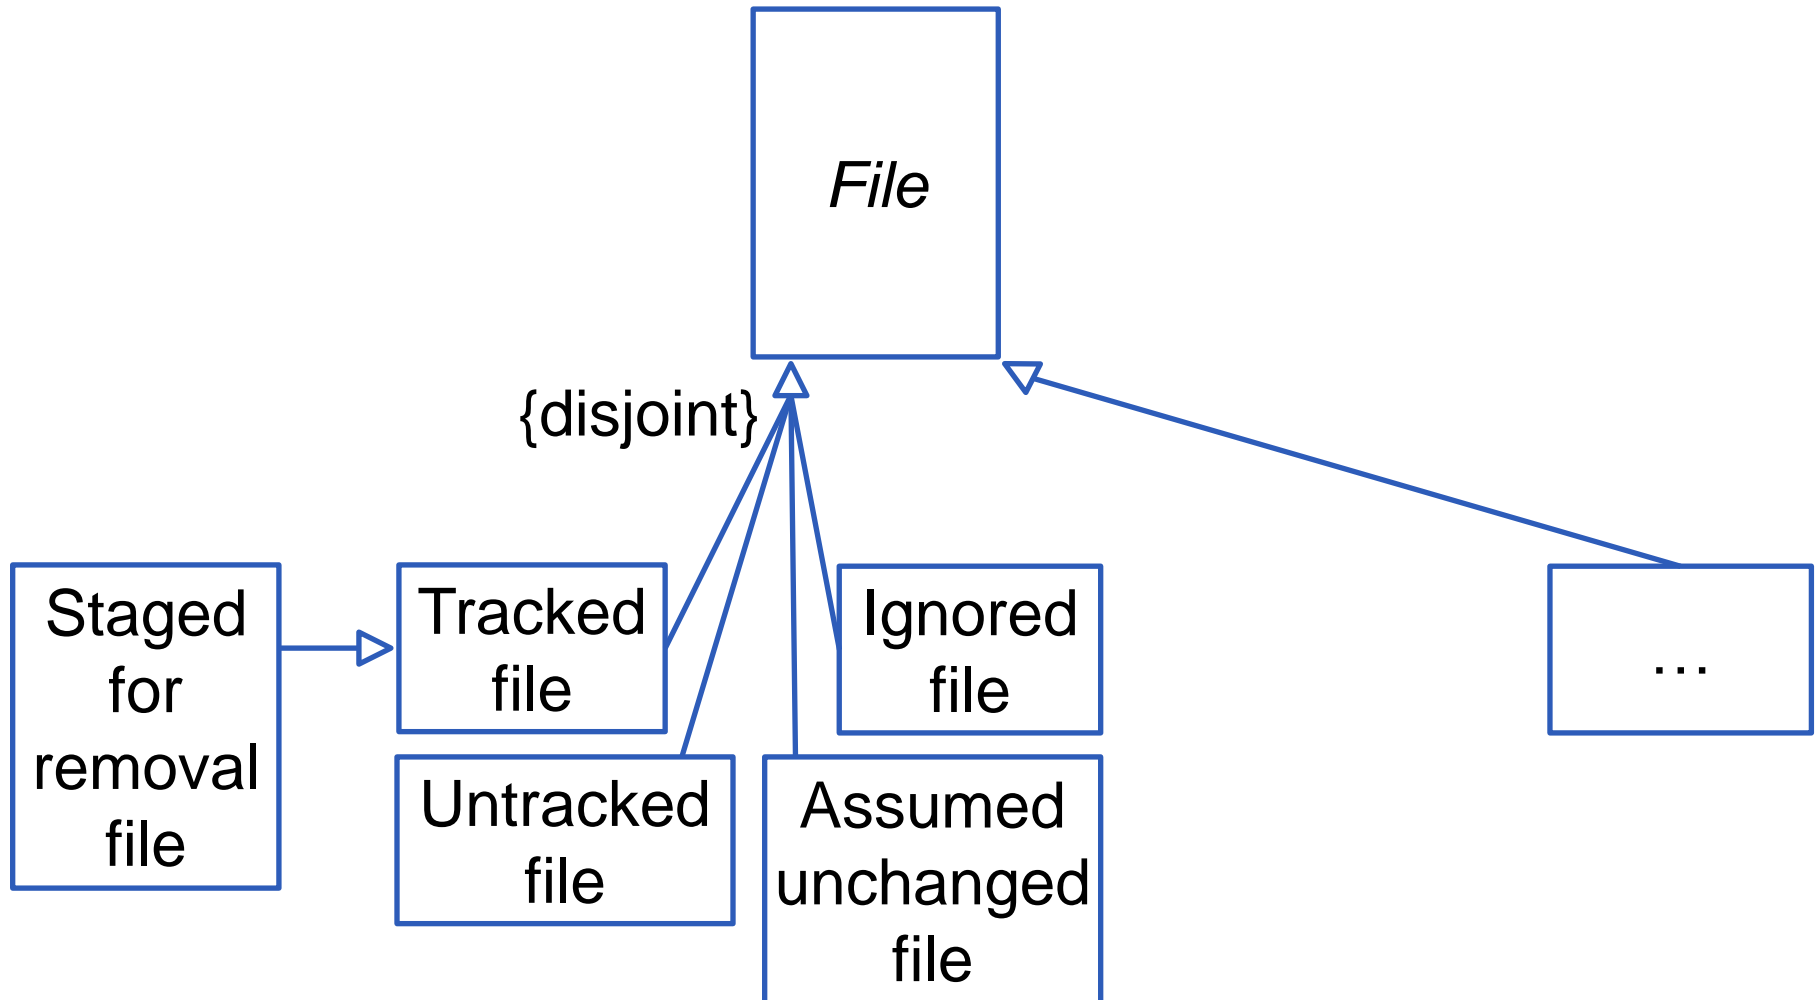
File italic = abstract
>> No other subtypes than the 5 shown here



(Un)Tracked → Assumed unchanged

- Tracked file
 - “File” whose modifications Git will notice
- Untracked file
 - “File” in the working directory & has no committed version
 - Adding such a file to Git → Git will start tracking it
- Assumed unchanged file
 - “File” that was previously tracked
 - User has indicated Git should no longer track it
 - Adding such a file to Git → No effect
 - Instead: Unset file’s “assume unchanged” bit

Conceptual model: Staged for removal

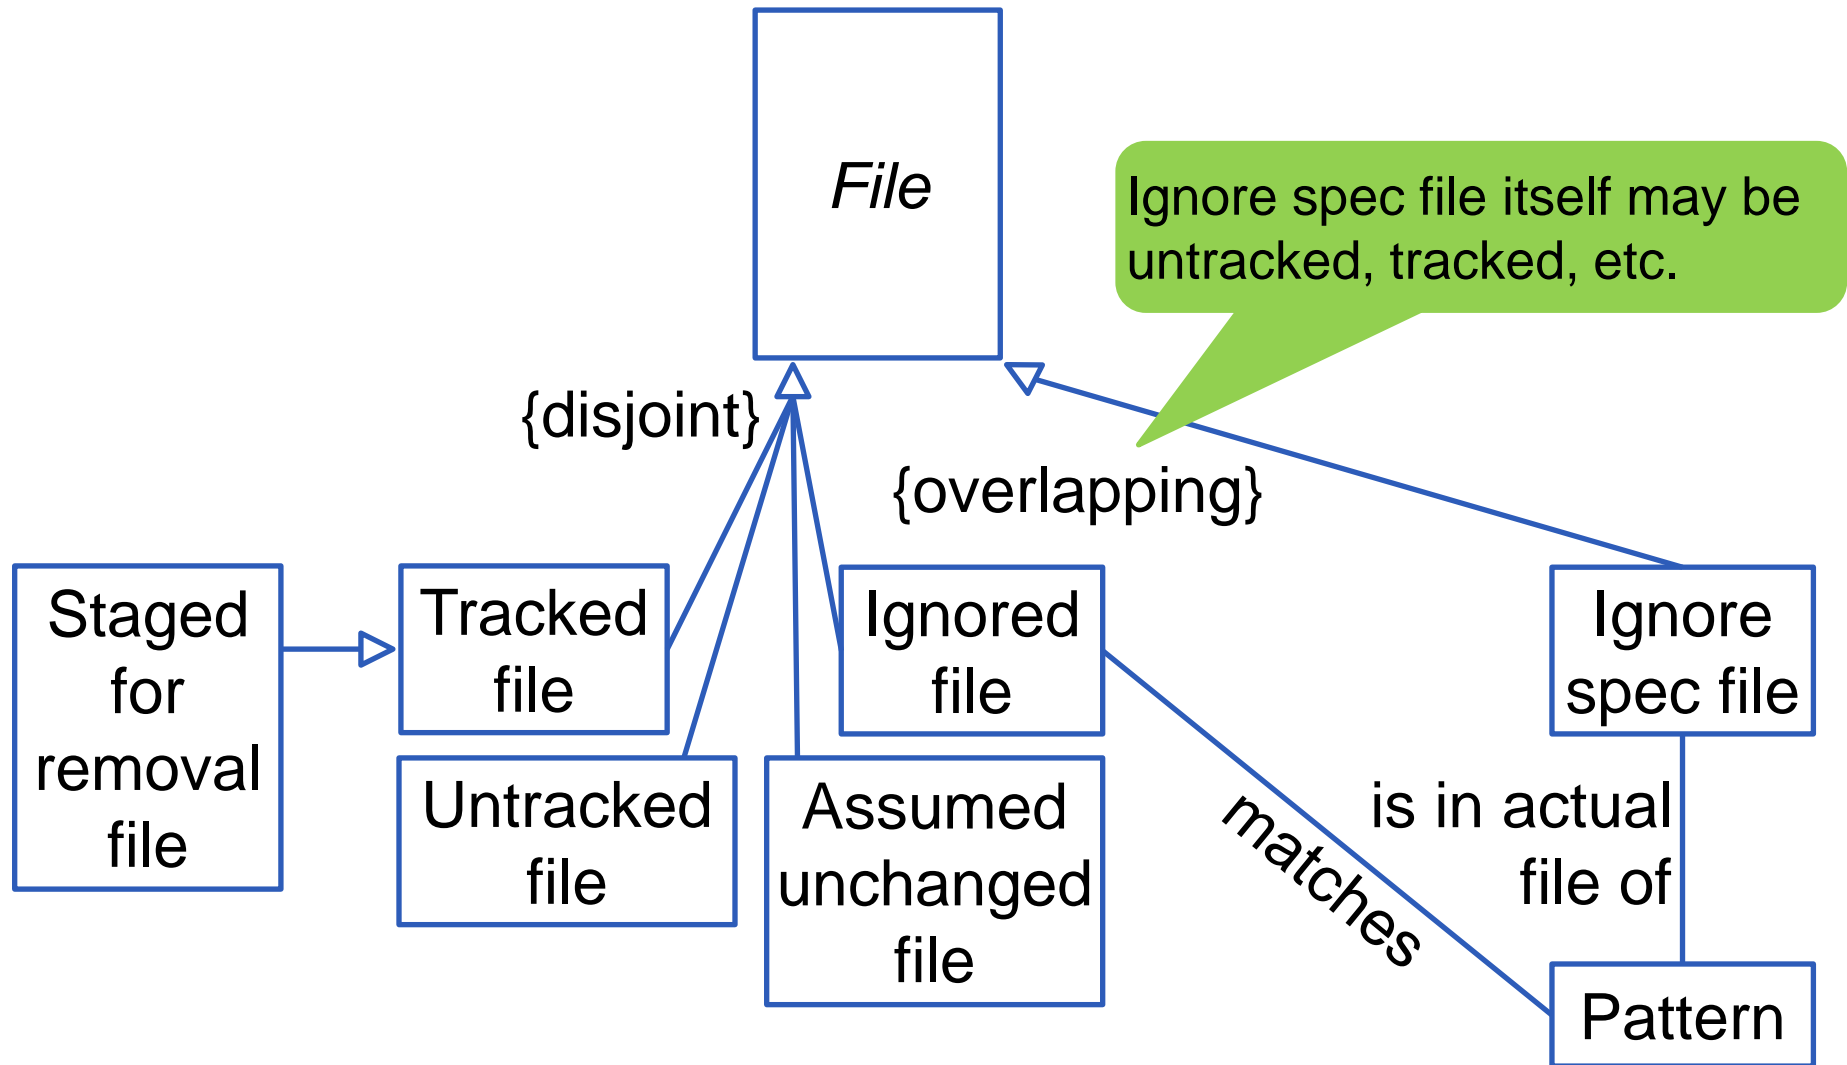


Based on [CD, Figure 1]

Staged for removal

- Staged for removal
 - File path of an actual file that no longer exists in working directory
 - Staging area records file's absence
 - Will be removed from the repo on next commit
- Staged for removal vs. un-staging
 - Un-staged := Staged version is removed

Conceptual model: .gitignore



Based on [CD, Figure 1]

Pattern

- String specifying a set of file paths
- Used in actual `.gitignore` files
- Example: `*.class` for Java project

Ignored

- Ignore spec file
 - File path of special kind of actual file (`.gitignore`)
 - Actual `.gitignore` file contains set of patterns of file paths Git should ignore
 - May be tracked, untracked, etc.
- Ignored file
 - File path ignored by Git
 - E.g.: Will not appear in output of “git status” command
 - Git will ignore “file” if it matches pattern in `.gitignore` in file’s directory or any of its (recursive) parent directories

Recall: Relation multiplicity constraint

- $A = \{a_1, a_2, \dots\}$

- $B = \{b_1, b_2, \dots\}$

- $\text{someRelationBetweenAandB} = \{(a_2, b_1), \dots\}$

Depending on the relation:

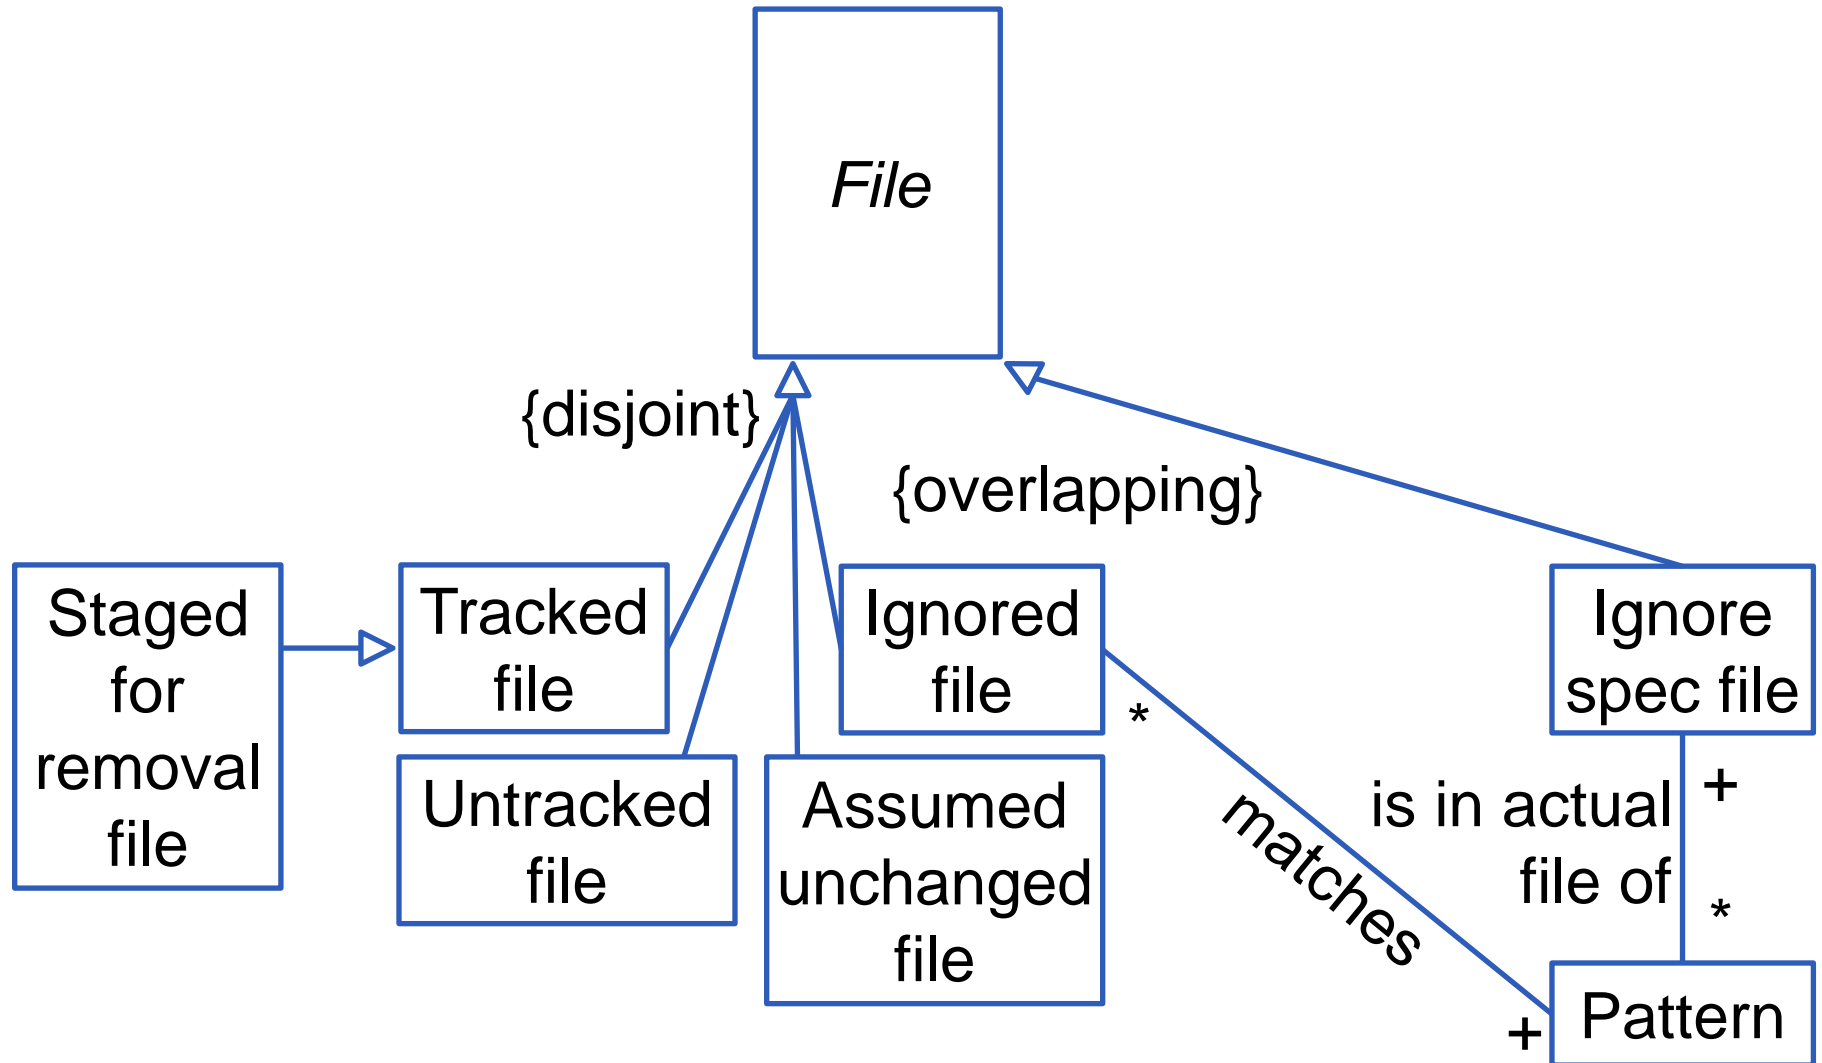
>> Some a_g, b_h may not appear in relation

>> Some a_i, b_j may appear several times



- Each item in A maps to z items in B
- Each item in B maps to x items in A
- Each annotation x, z may be a range or:
 - Star (*) = “ ≥ 0 ”, plus (+) = “ ≥ 1 ”
 - Question mark (?) = “0 or 1”

File (path) → Working, etc. version



Based on [CD, Figure 1]

Extra notation: Set of similar relations

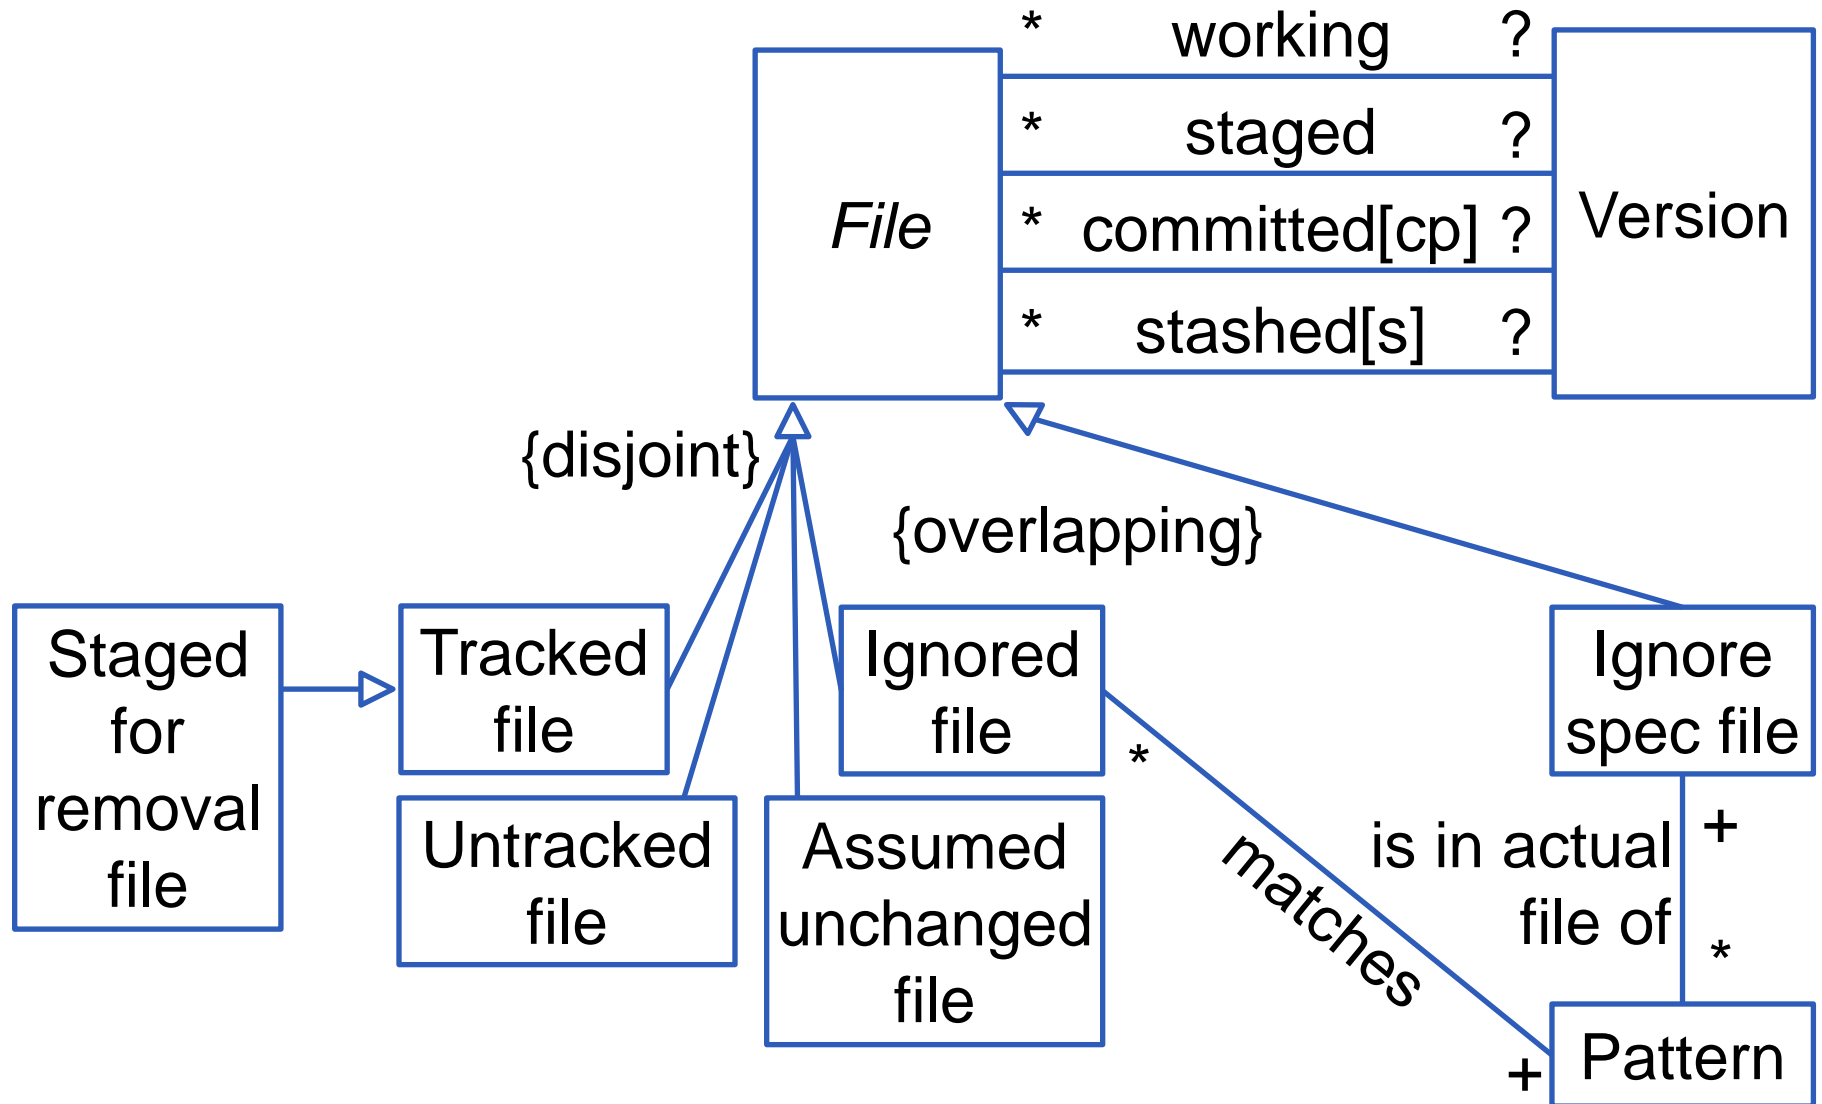
- $A = \{a_1, a_2, \dots\}$
- $B = \{b_1, b_2, \dots\}$
- $r[i] = \{(a_2, b_1), \dots\}$

>> “r” is a mapping from some “index” set to a set of relations
>> “i” is an element of that “index” set
>> Each “r[i]” is a relation between A and B



- In each $r[i]$: Each item in A maps to z items in B
- In each $r[i]$: Each item in B maps to x items in A
- Each annotation x, z may be a range or:
 - Star (*) = “ ≥ 0 ”, plus (+) = “ ≥ 1 ”
 - Question mark (?) = “0 or 1”

File (path) → Working, etc. version



Based on [CD, Figure 1]

Version, Branch, Stash

- Version := Contents of a file at some point in time
 - Two different files with same content → same version
- Version ≠ “version number”
 - Version number would be a name for a version
 - Git doesn't have concept of a version number
- Branch
 - Named collection of committed versions of files
 - Identifier, user can change name

Working version & Staged version

- Working
 - Working version of a file := stored in working directory
 - A file can have at most one working version
 - A file may be committed to repo but no longer be in working directory
- Staged
 - Staged version of a file := stored in staging area & saved to repo on commit
 - File can have at most one staged version
 - Untracked file cannot have a staged version
 - File can have a staged but no working version
 - If user deletes file from working directory after adding it to Git

Committed version & Stashed version

- Commit point (cp)
- Committed version
 - Committed version of a file := Stored in local repo at commit point cp
- Stash s := Collection of file versions saved together
 - Git maintains stack of all stashes
 - Stash name `stash{i}` assigned by Git (i = stack index)
- Stashed version
 - Stashed version of a file := Stored in stash s

(Implicit) git add → git commit

- `git add f`
 - Makes `f`'s working version also `f`'s staged version
 - If `f` is untracked → switch `f` to tracked
 - **Subsequent changes to actual file `f` will cause `f`'s working version to diverge from `f`'s staged version**
 - After changing contents of actual file `f` just run again: `git add f`
- `git commit`
 - Makes all **staged versions** also committed versions
- `git commit -a`
 - `git commit f1 f2`
 - Make **working versions** also committed versions
 - For all files or just for files `f1`, `f2`
 - Skips staging area (implicit `git add`)

Reverting

- `git reset HEAD f`
 - Disassociates `f` from staged version
 - If `f` was untracked before `git add` that staged it → switch `f` back to untracked
- `git checkout f`
 - Replace `f`'s working version with `f`'s staged version
 - No staged version → Replace working version with `f`'s committed version

Removing

- `rm f`
 - Remove f's working version (standard OS command)
 - Doesn't affect staged or committed versions
- `git rm f`
 - If f is tracked → stage f for removal, remove f from working directory
 - If f is assumed unchanged → same as above, but also f becomes tracked again

Following slides mostly based on Mike Ernst's best practices:

<http://homes.cs.washington.edu/~mernst/advice/version-control.html#best-practices>

BEST PRACTICES

Write a descriptive commit message

- Briefly explain purpose of this commit
 - In present tense
 - Refer to a specific issue this commit addresses
 - Include (link to) issue number in issue tracking system
- [-] Initially takes a bit of effort to write
 - But after some practice only takes a few seconds
- [+] Valuable later on
 - When searching in the history of commits

Each commit should refer to an issue

- First create an issue in your issue tracking system
- Now each commit should refer to an issue
- Easy in BitBucket, e.g.:
 - “#13: Fix license file”
 - In web view, BitBucket will link “#13” to issue 13
- [+] Makes it easier to
 - Understand individual commit
 - Find all commits for one issue
 - Since developers may work on multiple issues at the same time

Make each commit a logical unit

- Each commit should have a single purpose
- A commit should not serve two or more purposes
 - Makes it hard to undo one of these purposes
- When changing many files in working directory that address multiple purposes
 - Package these changes into separate commits
 - Do not have to commit all changes in single commit
- Each purpose should be contained in a single commit
- Do not spread a purpose over multiple commits
 - Makes it harder to undo a purpose (multiple commits)

Git workflow for creating a commit

- `git status`
 - Lists all the modified files
- `git diff`
 - Shows specific differences
 - Helps you compose a commit message
- `git commit f1 f2 -m "My commit message"`
 - Commits just the files f1, f2 you want to commit for this particular commit

Incorporate others' changes frequently

- Fetch others' commits daily, via one of:
 - `git pull`
 - `git pull -r`
- Reduces risk of merge conflicts

Frequently share your commits

- `git push`
- After completing each small task
- But avoid breaking the overall system
 - Check, verify, test before sharing your commits

Coordinate conflicting edits

- Ideally do not work on same file as your co-worker at the same time
 - Avoid later merge conflicts
- Person B should only start editing file F **after** person A committed and shared changes to F

Remember tools are line-based

- Tools look for changed lines and let users merge conflicting lines
- Keep each line short, ideally to about 74 characters
 - Fits on screen
 - Longer lines are harder to merge, have to scroll
- Changes are harder to spot in longer lines
 - Tool marks the entire line as changed
 - Often does not point to location in line

Do not commit generated files

- Do not add or commit, e.g.:
 - .class generated from your .java files
 - Exclude such files explicitly or via pattern
- User can re-generate such files from the sources
- Unnecessarily increases size of repository
- May trigger merge conflicts
 - Wastes users' time dealing with such conflicts
- But do add important binaries that rarely change
 - .jar of libraries needed to run code

Do not force things

- Using the -f force option can get you into bigger trouble